

# Exploratory Data Mining via Search Strategies Lab #3

Ross Jacobucci & Kevin J. Grimm

# Outline

This script will go over Decision Trees and generalizations in R.

Resources:

Basic intro to Decision Trees: <http://www.statmethods.net/advstats/cart.html>

Full list of data mining packages in R:

<http://cran.r-project.org/web/views/MachineLearning.html>

Two packages will be used and their caret equivalents:

- ▶ **rpart** (tree accomplishes very similar thing):<http://cran.r-project.org/web/packages/rpart/vignettes/longintro.pdf>
- ▶ **party**:  
<http://cran.r-project.org/web/packages/party/vignettes/party.pdf>

In **caret**, method =

- ▶ “rpart” – tuning = cp (complexity parameter)
- ▶ “rpart2” – tuning = maxdepth
- ▶ “rpartCost” – tuning = cp and cost
- ▶ “ctree” – tuning = mincriterion (p value thresholds)
- ▶ “ctree2” – tuning = maxdepth

# Lets load the main packages

```
library(caret)
library(rpart)
library(pROC)
library(randomForest)
library(ada)
library(ISLR)
library(party)
library(MASS) # for boston data
data(Boston)
```

# Regression (continuous outcome)

Use rpart first with the Boston data use regression first – predicting median value of homes


```
#str(Boston)  
  
# lets get a baseline with linear regression  
lm.Boston <- lm(medv ~., data=Boston)  
#summary(lm.Boston)
```

We do pretty well with linear regression R-squared of .74

# CART

How about if we just blindly apply Decision Trees

```
rpart.Boston <- rpart(medv ~., data=Boston)
#summary(rpart.Boston)
plot(rpart.Boston);text(rpart.Boston)
```



## CART Continued

```
pred1 <- predict(rpart.Boston)
cor(pred1, Boston$medv)**2
```

```
## [1] 0.8075721
```

Doing really well – Rsquared = 0.81

*# this can be hard to interpret, so I like to look at a different output*  
rpart.Boston

```
## n= 506
```

```
##
```

```
## node), split, n, deviance, yval
```

```
##      * denotes terminal node
```

```
##
```

```
## 1) root 506 42716.3000 22.53281
```

```
## 2) rm< 6.941 430 17317.3200 19.93372
```

```
## 4) lstat>=14.4 175 3373.2510 14.95600
```

```
## 8) crim>=6.99237 74 1085.9050 11.97838 *
```

```
## 9) crim< 6.99237 101 1150.5370 17.13762 *
```

```
## 5) lstat< 14.4 255 6632.2170 23.34980
```

```
## 10) dis>=1.5511 248 3658.3930 22.93629
```

```
## 20) rm< 6.543 193 1589.8140 21.65648 *
```

```
## 21) rm>=6.543 55 643.1691 27.42727 *
```

```
## 11) dis< 1.5511 7 1429.0200 38.00000 *
```

```
## 3) rm>=6.941 76 6059.4190 37.23816
```

```
## 6) rm< 7.437 46 1899.6120 32.11304
```

# Lasso Regression

What if we tried regularized (penalized) regression instead?

Note: for glmnet, both the x's and y have to be in separate matrices

- and all class = numeric

- don't worry about response, doesn't have to be factor for logistic

- just specify “binomial”

```
y.B <- Boston$medv
x.B <- sapply(Boston[, -14], as.numeric)

# alpha =1 for lasso, 0 for ridge
library(glmnet)
cv <- cv.glmnet(x.B, y.B, alpha=1)
lasso.reg <- glmnet(x.B, y.B, alpha=1, family="gaussian", lambda=cv$lambda.min)

lasso.resp <- predict(lasso.reg, newx=x.B)
cor(y.B, lasso.resp)**2
```

```
##                s0
## [1,] 0.7403028
```

Taking into account cross-validation, we do worse compared to linear regression with no tuning.


## CART Plots So the plot for rpart didn't come out that well.

Good news, there are better options for plotting.

# CART Plotting Continued

Note, `prp()` offers many additional capabilities for tweaking the plot For instance:

```
# ?prp  
prp(rpart.Boston,varlen=10,digits=5,fallen.leaves=T)
```






# CART Plotting Continued

Probably my favorite, `fancyRpartPlot()`

```
#fancyRpartPlot(); from rattle  
fancyRpartPlot(rpart.Boston)
```



Rattle 2016-Jul-26 18:16:05 R.Jacobucci

# Conditional Inference Trees

So what about with conditional inference trees?

What if we want a smaller tree? This can be accomplished a number of ways. We can prespecify the maxdepth, the minimum number of people per node, as well as making more restrictive splitting criterion.

Example of prespecifying the depth with ctree()

```
ctree.Boston <- ctree(medv ~., data=Boston)
#plot(ctree.Boston) # too big of a tree
pred2 <- predict(ctree.Boston)
cor(pred2,Boston$medv)**2
```

```
## [1] 0.8746338
```

We do better than rpart, Rsquared = 0.87

## Conditional Inference Trees Continued

Biggest difference between `ctree()` and `rpart()` is that `ctree()` does not demonstrate bias with respect to the number of response options, and supposedly had less of a propensity to overfit than `rpart()`.

Note: the models are not optimizing based on `Rsquared`, most likely `MSE`

So what do we think now? Are we happy with results? Remember, decision trees are generally quite robust, so it may not be necessary to check assumptions. – See Table 10.1 ESL

But what about generalizability?

Although not as serious as with SVM for instance, Decision Trees have a propensity to overfit, meaning the tree structure won't generalize well

So let's try just creating a simple Training and Test datasets

```
train = sample(dim(Boston)[1], dim(Boston)[1]/2) # half of sample
Boston.train = Boston[train, ]
Boston.test = Boston[-train, ]
```

# Linear Regression with CV

```
lm.train <- lm(medv ~., data=Boston.train)

pred.lmTest <- predict(lm.train,Boston.test)
cor(pred.lmTest,Boston.test$medv)**2
```

```
## [1] 0.7419161
```

Note: we are taking our lm object trained on the train dataset, and using these fixed coefficients to predict values on the test dataset.

In SEM, this is referred to as a tight replication strategy No difference in using a test dataset – both  $R^2$  are 0.74

How about with rpart?

# CART CV

```
rpart.train <- rpart(medv ~., data=Boston.train)

pred.rpartTest <- predict(rpart.train,Boston.test)
cor(pred.rpartTest,Boston.test$medv)**2
```

```
## [1] 0.7284012
```

Not as good – drops from 0.81 to 0.76 – still better than `lm()`


But with `rpart`, it is common to prune trees back. What if we try this, is there less of a drop in  $R^2$ ?

Note: `rpart` automatically does internal CV, varying the complexity parameter (`cp`). If you use the `tree` package instead, you will have to use `cv.tree()`

# CART Pruning

With `plotcp()` we are going to choose the error within 1 SE of the lowest cross-validated error. This will be used to prune

```
plotcp(rpart.train)
```



# CART Pruning Continued

```
printcp(rpart.train)

##
## Regression tree:
## rpart(formula = medv ~ ., data = Boston.train)
##
## Variables actually used in tree construction:
## [1] lstat nox    rm
##
## Root node error: 22730/253 = 89.84
##
## n= 253
##
##      CP nsplit rel error  xerror    xstd
## 1 0.463148      0   1.00000 1.00989 0.115443
## 2 0.157400      1   0.53685 0.66698 0.074911
## 3 0.106252      2   0.37945 0.49933 0.067967
## 4 0.033742      3   0.27320 0.35830 0.054779
## 5 0.021929      4   0.23946 0.32467 0.057639
## 6 0.012172      5   0.21753 0.28915 0.053367
## 7 0.011482      6   0.20536 0.28857 0.053482
## 8 0.010219      7   0.19387 0.29071 0.053492
## 9 0.010000      8   0.18366 0.28676 0.053575
```

What is xerror and other error?

# CART Pruning Continued


```
# rsq.rpart(rpart.train) ## another way to get cp plots
prune.Bos <- prune(rpart.train,0.053)

pred.prune <- predict(prune.Bos,Boston.test)
cor(pred.prune,Boston.test$medv)
```

```
## [1] 0.789153
```

Plot Pruned Tree

```
#plot(prune.Bos);text(prune.Bos)
fancyRpartPlot(prune.Bos)
```





```
ctree.train <- ctree(medv ~., data=Boston.train)
#plot(ctree.train) # huge tree
pred.ctreeTest <- predict(ctree.train,Boston.test)
cor(pred.ctreeTest,Boston.test$medv)**2
```

```
## [1] 0.732705
```

It is worth noting how much more of an effect there was for using a test dataset with the tree methods as compared to `lm()`, this is pretty typical, and much more important with more “flexible” methods such as random forests, `gbm`, `svm` etc. . .

# Classification

# Classification

## Two Biggest Things To Remember:


1. Make sure functions outcome variable is categorical; `as.factor(outcome)`
2. Using `predict()` changes. Variable across packages

As a baseline, we will use logistic regression.

```
library(ISLR)
data(Default)
#head(Default)
str(Default)
```

```
## 'data.frame':    10000 obs. of  4 variables:
## $ default: Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 1 1 1 ...
## $ student: Factor w/ 2 levels "No","Yes": 1 2 1 1 1 2 1 2 1 1 ...
## $ balance: num  730 817 1074 529 786 ...
## $ income : num  44362 12106 31767 35704 38463 ...
```

My favorite function in R is `str()`, as it gives the class of each variable and other summary characteristics. Most important thing to note is that the “default” variable is already coded as a factor variable, meaning that R now knows it is categorical, and will change the cost function (thus estimator) accordingly.

This is really important because `rpart`, `randomForest` and other packages do not automatically detect whether it is a regression or classification problem. If you don't change the outcome variable to its proper class, you could get a suboptimal answer (use 

# ROC Curves

good intro to using ROC:

<https://ccrma.stanford.edu/workshops/mir2009/references/ROCintro.pdf>

These plots are a balance of sensitivity and specificity. Ideally the curve gets as close as possible to the upper left corner.

To get this plot, we need to get our predictions from our logistic model.

```
glm.probs=predict(lr.out,type="response")
#glm.pred00=ifelse(glm.probs>0.5,1,0)

rocCurve <- roc(Default$default,glm.probs)
pROC::auc(rocCurve)
```


```
## Area under the curve: 0.9496
```

```
pROC::ci(rocCurve)
```

```
## 95% CI: 0.9402-0.959 (DeLong)
```

# Plot ROC Curve

```
# quartz()  
plot(rocCurve, legacy.axes = TRUE, print.thres=T, print.auc=T)
```



# Lasso Logistic Regression

How about lasso logistic regression?

```
library(glmnet)
yy = as.numeric(Default$default)
xx = apply(Default[,2:4],as.numeric)
lasso.out <- cv.glmnet(xx,yy,family="binomial",alpha=1,nfolds=10) #alpha=1 ==
# find best lambda
ll <- lasso.out$lambda.min

lasso.probs <- predict(lasso.out,newx=xx,s=ll,type="response")
```

Results from lasso using CV

```
rocCurve.lasso <- roc(Default$default,lasso.probs)
pROC::auc(rocCurve.lasso)
```


## Area under the curve: 0.9495

```
pROC::ci(rocCurve.lasso)
```

## 95% CI: 0.9401-0.959 (DeLong)

# Plot ROC


```
# quartz()  
plot(rocCurve.lasso, legacy.axes = TRUE, print.thres=T, print.auc=T)
```



# Using Decision Trees for Classification

Instead of demonstrating how to use `rpart()` or `ctree()`, I prefer to use the `train()` from `caret`. This makes it much easier to test out multiple different methods, as well as automatically vary the tuning parameters such as depth, complexity etc..  
`train()` for `ctree`

```
train.ctree <- train(as.factor(default)~student+balance+income,data=Default,method='ctree',  
plot=train.ctree)
```





## using a confusion matrix

```
train.class <- predict(train.rpart,Default,type="raw")
confusionMatrix(train.class,Default$default)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    No  Yes
```

```
##           No 9639 243
```

```
##           Yes   28   90
```

```
##
```

```
##                Accuracy : 0.9729
```

```
##                95% CI : (0.9695, 0.976)
```

```
##      No Information Rate : 0.9667
```

```
##      P-Value [Acc > NIR] : 0.0002082
```

```
##
```

```
##                Kappa : 0.3885
```

```
##  Mcnemar's Test P-Value : < 2.2e-16
```

```
##
```

```
##                Sensitivity : 0.9971
```

```
##                Specificity : 0.2703
```

```
##      Pos Pred Value : 0.9754
```

```
##      Neg Pred Value : 0.7627
```

```
##      Prevalence : 0.9667
```

```
##      Detection Rate : 0.9639
```

```
##      Detection Prevalence : 0.9882
```

```
##      Balanced Accuracy : 0.6337
```

# table()

The typical way of getting a table of classification results:

```
table(train.class,Default$default)
```

```
##  
## train.class    No   Yes  
##           No 9639  243  
##           Yes  28   90
```

I use `confusionMatrix()` because I am lazy and don't want to calculate all of the other statistics

## Changing the cutoff for class assignment

This uses the optimal cutoff from the pROC plot

```
train.prob <- predict(train.rpart,Default,type="prob")[,2]
train.class2 <- ifelse(train.prob > .031,1,0)
confusionMatrix(as.numeric(Default$default)-1,train.class2,positive="1")
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
## Prediction    0    1
##           0 9541 126
##           1  171 162
```

```
##
```

```
##           Accuracy : 0.9703
##           95% CI : (0.9668, 0.9735)
##           No Information Rate : 0.9712
##           P-Value [Acc > NIR] : 0.71715
```

```
##
```

```
##           Kappa : 0.5065
##           Mcnemar's Test P-Value : 0.01068
```


```
##
```

```
##           Sensitivity : 0.5625
##           Specificity : 0.9824
##           Pos Pred Value : 0.4865
##           Neg Pred Value : 0.9870
##           Prevalence : 0.0288
```

```
##
```

## Another Way to Get Optimal Threshold

```
library(ROCR)
pred <- prediction(train.prob,as.numeric(Default$default)-1)
perf <- performance(pred,"tpr","fpr")
#str(perf)
plot(perf)
```



```
cutoffs <- data.frame(cut=perf@alpha.values[[1]], fpr=perf@x.values[[1]],
                      tpr=perf@y.values[[1]])
#cutoffs
```